

pp4fpga Sparse Matrix Vector Multiplication report

電子所碩一 R09943019 李唐

➤ System introduction

Sparse matrix vector (SpMV) multiplication performs the operation

$$y = M * x$$

where x and y are vectors, and M is a sparse matrix stored in the compressed row storage (CRS) form.

CRS form is commonly used when the matrix contains only a small number of nonzero elements (typically 10 percent or less), enabling a matrix to be stored with less memory and manipulated with fewer operations.

$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \end{bmatrix} = \begin{bmatrix} m_{11} & m_{12} & 0 & 0 \\ 0 & m_{22} & m_{23} & 0 \\ m_{31} & 0 & m_{33} & m_{34} \\ 0 & m_{42} & 0 & m_{44} \end{bmatrix} * \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix}$$

Given a sparse matrix M in 2-dim representation as shown above, the corresponding CRS representation consists of three arrays:

✧ **values**

The values array holds the value of each nonzero element in the matrix.

$$[m_{11} \quad m_{12} \quad m_{22} \quad m_{23} \quad m_{31} \quad m_{33} \quad m_{34} \quad m_{42} \quad m_{44}]$$

✧ **columnIndex**

The columnIndex array stores the column index of each nonzero element.

$$[0 \quad 1 \quad 1 \quad 2 \quad 0 \quad 2 \quad 3 \quad 1 \quad 3]$$

✧ **rowPtr**

The rowPtr array contains the index in values array of the first element in each row. In this example, the index of m_{11} , m_{22} , m_{31} , m_{42} , and the total number of nonzero elements in values array are stored in rowPtr.

$$[0 \quad 2 \quad 4 \quad 7 \quad 9]$$

※ If there are rows in the matrix without a nonzero element, then values in the rowPtr array will repeat.

Given the CRS form representation, we can conclude that

- if $\text{values}[k] = m_{ij}$, then $\text{colIndex}[k] = j$
- if $\text{values}[k] = m_{ij}$, then $\text{rowPtr}[i] \leq k < \text{rowPtr}[i + 1]$
- the number of nonzero elements in row i is $\text{rowPtr}[i + 1] - \text{rowPtr}[i]$

Sparse matrix vector multiplication can be more efficiently computed by representing the matrix in CRS form, which can significantly reduce the operations required especially when the matrix is highly sparse.

➤ Screen dump and Observations

✧ **Solution 1: Baseline implementation**

```
L1: for (int i = 0; i < NUM_ROWS; i++) {
    y0 = 0;
    L2: for (int k = rowPtr[i]; k < rowPtr[i+1]; k++) {
        y0 += values[k] * x[columnIndex[k]];
    }
    y[i] = y0;
}
```

L1 iterates through each row of the matrix, and L2 iterates across the column of the matrix to accumulate the partial sum.

✧ Synthesis report/RTL Co-simulation

Performance Estimates

▢ Timing

▢ Summary

Clock	Target	Estimated	Uncertainty
ap_clk	10.00 ns	8.024 ns	1.25 ns

▢ Latency

▢ Summary

Latency (cycles)		Latency (absolute)		Interval (cycles)		
min	max	min	max	min	max	Type
13	189	0.130 us	1.890 us	13	189	none

▢ Detail

⊕ Instance

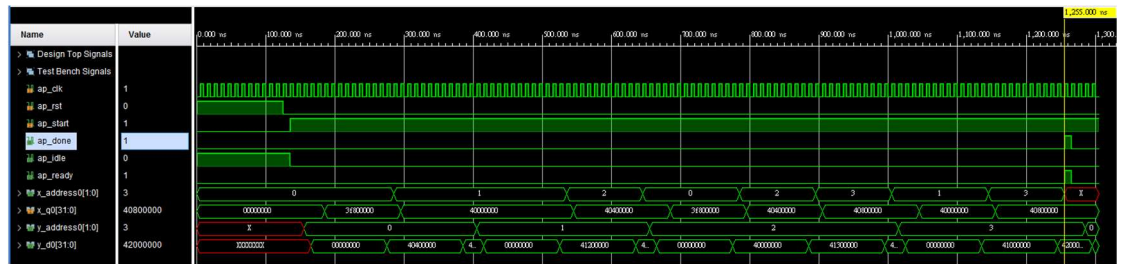
▢ Loop

Loop Name	Latency (cycles)		Iteration Latency	Initiation Interval		Trip Count	Pipelined
	min	max		achieved	target		
- L1	12	188	3 ~ 47	-	-	4	no
+ L2	0	44	11	-	-	0 ~ 4	no

Utilization Estimates

▢ Summary

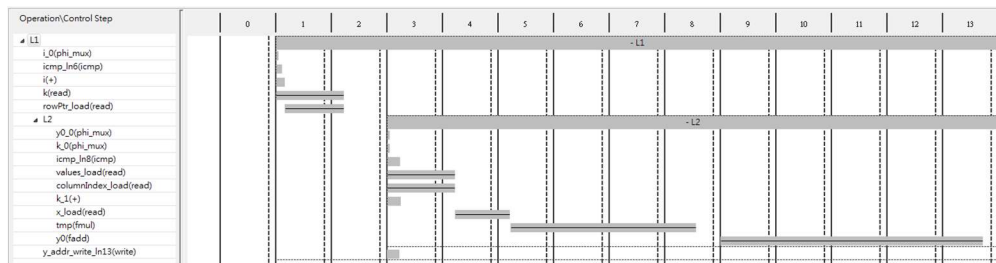
Name	BRAM_18K	DSP48E	FF	LUT	URAM
DSP	-	-	-	-	-
Expression	-	-	0	78	-
FIFO	-	-	-	-	-
Instance	-	5	348	711	-
Memory	-	-	-	-	-
Multiplexer	-	-	-	89	-
Register	-	-	215	-	-
Total	0	5	563	878	0
Available	280	220	106400	53200	0
Utilization (%)	0	2	~0	1	0



Total latency (cycles)	L2 max iteration latency (cycles)
112	36

The sparse matrix-vector multiplication requires variable operations in each iteration. Thus, the latency estimated in synthesis report is not accurate. Instead, we measure the latency from the cosim waveform.

✧ Analysis perspective



	Pipelined	Latency	Iteration Latency	Initiation Interval	Trip count
● spmv	-	13~189	-	14 ~ 190	-
● L1	no	12 ~ 188	3 ~ 47	-	4
●	no	0 ~ 44	11	-	0~4

✧ Solution 2: Pipeline L1

```

L1: for (int i = 0; i < NUM_ROWS; i++) {
# pragma HLS PIPELINE
    y0 = 0;
    L2: for (int k = rowPtr[i]; k < rowPtr[i+1]; k++) {
        y0 += values[k] * x[columnIndex[k]];
    }
    y[i] = y0;
}

```

✧ Synthesis report/RTL Co-simulation

Performance Estimates

Timing

Summary

Clock	Target	Estimated	Uncertainty
ap_clk	10.00 ns	8.024 ns	1.25 ns

Latency

Summary

Latency (cycles)		Latency (absolute)		Interval (cycles)		
min	max	min	max	min	max	Type
13	189	0.130 us	1.890 us	13	189	none

Detail

Instance

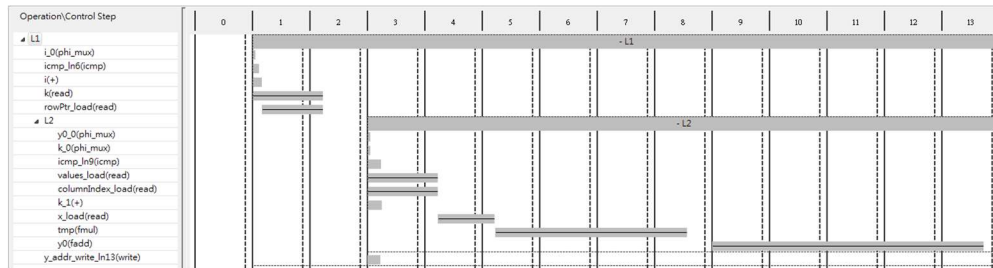
Loop

Loop Name	Latency (cycles)		Iteration Latency	Initiation Interval		Trip Count	Pipelined
	min	max		achieved	target		
- L1	12	188	3 ~ 47	-	-	4	no
+ L2	0	44	11	-	-	0 ~ 4	no

Total latency (cycles)	L2 max iteration latency (cycles)
112	36

There is no change to the overall performance since the L2 loop iteration is variable bounded. Thus, the L2 loop cannot be flattened.

✧ Analysis perspective



	Pipelined	Latency	Iteration Latency	Initiation Interval	Trip count
spmv	-	13~189	-	14 ~ 190	-
L1	no	12 ~ 188	3 ~ 47	-	4
	no	0 ~ 44	11	-	0~4

✧ Solution 3: Pipeline L2

```

L1: for (int i = 0; i < NUM_ROWS; i++) {
    y0 = 0;
    L2: for (int k = rowPtr[i]; k < rowPtr[i+1]; k++) {
# pragma HLS PIPELINE
        y0 += values[k] * x[columnIndex[k]];
    }
    y[i] = y0;
}

```

✧ Synthesis report/RTL Co-simulation

Performance Estimates

Timing

Summary

Clock	Target	Estimated	Uncertainty
ap_clk	10.00 ns	8.024 ns	1.25 ns

Latency

Summary

Latency (cycles)		Latency (absolute)		Interval (cycles)		
min	max	min	max	min	max	Type
17	117	0.170 us	1.170 us	17	117	none

Detail

Instance

Loop

Loop Name	Latency (cycles)		Iteration Latency	Initiation Interval		Trip Count	Pipelined
	min	max		achieved	target		
- L1	16	116	4 ~ 29	-	-	4	no
+ L2	0	25	11	5	1	0 ~ 4	yes

Utilization Estimates

Summary

Name	BRAM_18K	DSP48E	FF	LUT	URAM
DSP	-	-	-	-	-
Expression	-	-	0	80	-
FIFO	-	-	-	-	-
Instance	-	5	348	711	-
Memory	-	-	-	-	-
Multiplexer	-	-	-	92	-
Register	-	-	215	-	-
Total	0	5	563	883	0
Available	280	220	106400	53200	0
Utilization (%)	0	2	~0	1	0

Instance


Instance	Module	BRAM_18K	DSP48E	FF	LUT	URAM
spmv_fadd_32ns_32bkb_U1	spmv_fadd_32ns_32bkb	0	2	205	390	0
spmv_fmud_32ns_32cud_U2	spmv_fmud_32ns_32cud	0	3	143	321	0
Total		2	5	348	711	0

Interface

Summary

RTL Ports	Dir	Bits	Protocol	Source Object	C Type
ap_clk	in	1	ap_ctrl_hs	spmv	return value
ap_rst	in	1	ap_ctrl_hs	spmv	return value
ap_start	in	1	ap_ctrl_hs	spmv	return value
ap_done	out	1	ap_ctrl_hs	spmv	return value
ap_idle	out	1	ap_ctrl_hs	spmv	return value
ap_ready	out	1	ap_ctrl_hs	spmv	return value
rowPtr_address0	out	3	ap_memory	rowPtr	array
rowPtr_ce0	out	1	ap_memory	rowPtr	array
rowPtr_q0	in	32	ap_memory	rowPtr	array
rowPtr_address1	out	3	ap_memory	rowPtr	array
rowPtr_ce1	out	1	ap_memory	rowPtr	array
rowPtr_q1	in	32	ap_memory	rowPtr	array
columnIndex_address0	out	4	ap_memory	columnIndex	array
columnIndex_ce0	out	1	ap_memory	columnIndex	array
columnIndex_q0	in	32	ap_memory	columnIndex	array

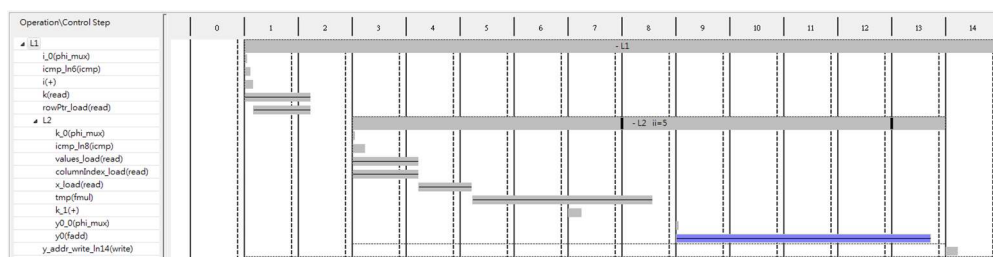
values_address0	out	4	ap_memory	values	array
values_ce0	out	1	ap_memory	values	array
values_q0	in	32	ap_memory	values	array
y_address0	out	2	ap_memory	y	array
y_ce0	out	1	ap_memory	y	array
y_we0	out	1	ap_memory	y	array
y_d0	out	32	ap_memory	y	array
x_address0	out	2	ap_memory	x	array
x_ce0	out	1	ap_memory	x	array
x_q0	in	32	ap_memory	x	array



Total latency (cycles)	L2 max iteration latency (cycles)
82	24

Pipelining to the inner L2 loop can reduce the latency required. However, the II is limited by the floating-point addition in each iteration.

✧ Analysis perspective



	Pipelined	Latency	Iteration Latency	Initiation Interval	Trip count
● spmv	-	17~117	-	18 ~ 118	-
● L1	no	16 ~ 116	4 ~ 29	-	4
●	yes	0 ~ 25	11	5	0~4

The blue mark shows the limiting operation.

✧ **Solution 4: Pipeline L2 + Unroll L2 with factor 4**

```

L1: for (int i = 0; i < NUM_ROWS; i++) {
    y0 = 0;
    L2: for (int k = rowPtr[i]; k < rowPtr[i+1]; k++) {
# pragma HLS PIPELINE
# pragma HLS UNROLL factor=4
        y0 += values[k] * x[columnIndex[k]];
    }
    y[i] = y0;
}

```

✧ Synthesis report/RTL Co-simulation

Performance Estimates

Timing

Summary

Clock	Target	Estimated	Uncertainty
ap_clk	10.00 ns	8.024 ns	1.25 ns

Latency

Summary

Latency (cycles)		Latency (absolute)		Interval (cycles)		
min	max	min	max	min	max	Type
101	181	1.010 us	1.810 us	101	181	none

Detail

Instance

Loop

	Latency (cycles)			Initiation Interval			
Loop Name	min	max	Iteration Latency	achieved	target	Trip Count	Pipelined
- L1	100	180	25 ~ 45	-	-	4	no
+ L2	21	41	26	20	1	0 ~ 1	yes

Utilization Estimates


Summary

Name	BRAM_18K	DSP48E	FF	LUT	URAM
DSP	-	-	-	-	-
Expression	-	-	0	253	-
FIFO	-	-	-	-	-
Instance	-	5	348	711	-
Memory	-	-	-	-	-
Multiplexer	-	-	-	381	-
Register	-	-	651	-	-
Total	0	5	999	1345	0
Available	280	220	106400	53200	0
Utilization (%)	0	2	~0	2	0

Total latency (cycles)	L2 max iteration latency (cycles)
101	37

Unrolling the L2 loop does no help in latency since the floating point addition is performed successively.

✧ Analysis perspective



	Pipelined	Latency	Iteration Latency	Initiation Interval	Trip count
spmv	-	101~181	-	102 ~ 182	-
L1	no	100 ~ 180	25 ~ 45	-	4
	yes	21 ~ 41	26	20	0~1

➤ GitHub submission

https://github.com/nodetibyln/MSoC_self_paced_learning

➤ Reference

1. R. Kastner, J. Matai, and S. Neuendorffer , Parallel Programming for FPGAs, arXiv , 2018.